

Python Interview Prep — 2026

Edition · 00 · START HERE

Which file should I open?

You just opened a Drive folder with 10 PDFs. Your interview is somewhere between five days and a month away, and you're staring at filenames. **Don't open any of them yet.** Spend the next 60 seconds on this file instead. It exists for exactly one reason — to tell you which of the other nine PDFs to open right now, in what order, and what to skip. This pack assumes you already write Python; nothing in here re-teaches `for` loops or what a list is. The angle is simple: writing Python every day is not the same skill as clearing a Python interview, and the gap is what catches people. Read the diagnostic, find your row, close this file, go.

The 3-Question Diagnostic

Answer these three in your head. Then jump to the routing plans below.

Q1. When is your interview?

- (A) This week — 3 to 7 days out
- (B) 1 to 3 weeks out
- (C) 1 month or more
- (D) No specific date — building the skill

Q2. What role are you targeting?

- (i) Data Analyst
- (ii) Data Scientist
- (iii) Data Engineer
- (iv) ML Engineer
- (v) Backend Developer (Django / Flask / FastAPI)
- (vi) Full-stack Developer (Python + JS frontend)
- (vii) SDE — generalist (interview will cover both data and backend flavours)

Q3. Seniority?

- (1) Fresher — final year or 0-1 yr
- (2) 1 to 3 years
- (3) 3 to 7 years
- (4) 7 years and above

Roles (i) through (iv) are your **Data track**. Roles (v) and (vi) are your **Backend track**. Role (vii) is the **generalist** case — you read both tracks, lighter. This single fork drives most of what's in the routing plans below; get it wrong and you'll waste 40% of your prep time on questions your panel won't ask. The rest of this file assumes you've committed to a track.

Track Routing — Pick One Before You Plan

Your role	Heavy file	Skip entirely	Why
Data Analyst · Data Scientist · Data Engineer · ML Engineer	File 06 (Data Track Library — 200 Qs · pandas, numpy, wrangling)	File 07 (Backend Track)	Your panel is testing pandas / numpy / SQL-thinking-in-Python. Django and FastAPI questions don't show up in data-role interviews.
Backend Developer · Full-stack Developer	File 07 (Backend Track Library — 150 Qs · DSA, OOP, Django, Flask, FastAPI, REST, DB)	File 06 (Data Track)	Your panel is testing DSA-in-idiomatic-Python plus a framework round. <code>pd.merge</code> semantics don't come up; <code>select_related</code> does.
SDE generalist	Both 06 + 07 , lighter pass	None — but plan double the time	The role lands somewhere between the two and the panel hedges. Slightly heavier on Backend because round 2 is usually framework-flavoured even in generalist tracks.

Special-case note for ML Engineers: if the role is heavy on production ML serving (FastAPI + model endpoints + Redis caching), treat yourself as a generalist and add File 07 Chapter 4 (Flask + FastAPI) and Chapter 5 (REST API design) to your plan. Most ML roles in Indian product cos and GCCs are not pure research — they ship endpoints.

Everyone, regardless of track, reads **File 04 (Core Python Library)**, **File 05 (Guess the Output — 50 Traps)**, and **File 08 (Cheatsheet)**. Those three are the language itself; the track files are just where it lands on your interview day.

⚡ The 5-Day Cram Plan

You have 3 to 7 days. Total commitment: ~3 hours/day weekdays, 6 hours on one weekend day. About 18 hours of focused study.

The order matters. Traps first because they're the fastest confidence-builder and they prime you to spot the gotchas. Then Top 50 to compress the surface area. Then your track's top 50 because that's what round 1 will actually test. Cheatsheet the night before so it's fresh.

Day 1 (3 hrs) — Traps and the Top 50 - Hour 1–2: **File 05** — work the first 25 traps. Read the snippet, predict the output, then read the reveal and the *why*. Do not skim the "why" — the why is what you say to the interviewer. - Hour 2–3: **File 05** — the remaining 25 traps. By the end you should know the eight categories cold: mutable defaults, late-binding closures, `is` vs `==`, integer caching, float math, list aliasing, comprehension scoping, iterator exhaustion. - Hour 3–4: **File 01 Top 50** — read straight through. Don't try to solve. Mark every question you couldn't have answered cold.

Day 2 (3 hrs) — Core concepts that always show up - Hour 1–2: From **File 04**, Chapter 2 (Decorators) — first 20 Qs. The "what runs first" with stacked decorators is the single most-asked decorator question in Indian interviews. - Hour 2–3: From **File 04**, Chapter 3 (Generators + iterators) — first 20 Qs. Cover the memory question (`yield` vs returning a list) until you can say it in 30 seconds. - Hour 3–4: From **File 04**, Chapter 4 (The GIL + concurrency) — first 15 Qs. The GIL is interview catnip; you need a 60-second explanation that doesn't lie.

Day 3 (3 hrs) — Your track, top 50 - Hour 1–2: Open your track file (**06** if Data, **07** if Backend). Read the chapter table of contents. Then work the first chapter end to end. - Hour 2–3: Continue with chapter 2 of your track. (For Data: groupby + agg. For Backend: DSA in idiomatic Python.) - Hour 3–4: Cover chapter 3 of your track. (For Data: merge / join / concat. For Backend: OOP design plus your framework — Django or FastAPI depending on the role.)

Day 4 (3 hrs) — OOP, mutability, what hurts most candidates - Hour 1–2: From **File 04**, Chapter 1 (OOP + dunder) — first 25 Qs. The `__init__` vs `__new__` distinction and MRO with diamond inheritance are the two failures we see most. - Hour 2–3: From **File 04**, Chapter 5 (Memory + mutability) — first 20 Qs. `id()` vs `is` vs `==`, shallow vs deep copy, why `a = b = []` bites people. - Hour 3–4: Re-attempt every question you flagged on Days 1–3. Speak the answer out loud — interviews are spoken, not written.

Day 5 (6 hrs, weekend) — Mock pass and cheatsheet drill - Hour 1–2: Re-read **File 01 Top 50**. You should now answer 40+ of 50 cold. Mark gaps. - Hour 3–4: Speed-run your track's remaining chapters from **File 06** or **07**. Read the question, give yourself 90 seconds to think, then read the answer. Don't write code. - Hour 5: Re-attempt all gaps from the week. Speak each answer out loud, timed at 60 seconds per question. - Hour 6: **File 08 Cheatsheet** — read it three times. Cover the page, recite. Print it.

What you skip on a 5-day plan: **File 02** full read, **File 03**, the off-track file (06 or 07 — whichever your role doesn't need), **File 09**, the back half of **File 04**'s tail chapters (typing, modules, strings & regex unless you're a generalist).

Night before the interview: re-read **File 08** Cheatsheet once. Skim the eight trap categories from **File 05**. Sleep early. Eat properly. Do not study new material — it shakes confidence, never builds it.

The 14-Day Prep Plan

You have 2 to 3 weeks. ~2 hours/day weekdays, 4 hours each weekend day. About 36 hours.

Everything in the 5-day plan, **plus** the breadth that separates "knew enough" from "walked in prepared."

Days 1–5: run the 5-day plan above at slightly slower pace (2 hours weekday cap; let yourself absorb).

Days 6–8 — File 02 Top 200, in tier order - Day 6: All 50 **S-tier** Qs from **File 02** (these overlap File 01 — fast pass). - Day 7: First 45 of 90 **A-tier** Qs from **File 02**. These are where the decorator edge cases, dunder methods, and concurrency depth live. - Day 8: Remaining 45 A-tier Qs.

Days 9–10 — File 04 Core, by your weakness - Pick your two weakest chapters from File 04 based on what you flagged in week 1. Common weaknesses for Indian buyers, in order of how often they bite: Chapter 4 (GIL + concurrency), Chapter 7 (Comprehensions + scoping), Chapter 8 (Closures + late-binding), Chapter 9 (Typing + Protocol). - Work each chapter end to end. Speak the 30-second answer out loud per Q.

Days 11–12 — Your track, deeper - Continue your track file. Aim for 80% coverage of File 06 (Data) or File 07 (Backend) by end of day 12. - For Data: prioritise apply / vectorise (Ch 5) and the "here's a CSV" scenarios (Ch 8) — these are round 2 material. - For Backend: prioritise REST API design (Ch 5), DB interaction with N+1 framing (Ch 6), and the system scenarios (Ch 7).

Day 13 — File 02 B-tier + File 05 re-pass - 60 B-tier Qs from **File 02** — descriptors, metaclasses, ABCs, weakref. You will likely not be asked these, but for product-co rounds at Razorpay / Postman / Zerodha and GCC rounds at Goldman / JPMC, one or two might appear and you cannot freeze. - Re-run all 50 traps from **File 05**. By now you should predict the output before reading the answer on 40+ of 50.

Day 14 (interview eve) — File 08 + File 01 skim - Read **File 08** Cheatsheet. Memorise the 8 traps row and the time-complexity table. - Skim **File 01** to confirm you can speak any Q in 60 seconds. - Stop. Sleep early.

The 30-Day Deep Prep Plan

You have a month or more. ~1 hour/day weekdays, 3 hours each weekend day. About 46 hours.

The difference between this and the 14-day plan: you actually finish File 04, you do the cutting-edge intel from File 03, and you do daily warmups with File 05 instead of cramming the traps.

Daily warmup (every day, 15 min): 2 traps from File 05. Cover the answer. Predict. Reveal. Read the why aloud. By day 25 you'll have cycled all 50 traps twice. This is the cheapest, highest-leverage habit in the whole plan.

Week 1 — Core foundations - Days 1–2: **File 01 Top 50** end to end, written answers where useful. - Days 3–5: **File 04** Chapters 1 (OOP + dunder) and 2 (Decorators). Slow reading. - Weekend: **File 04** Chapter 3 (Generators + iterators) + Chapter 5 (Memory + mutability).

Week 2 — The concepts that decide rounds - Days 6–8: **File 04** Chapter 4 (GIL + concurrency). Single hardest chapter in the pack for most buyers; pace it. - Days 9–10: **File 04** Chapter 6 (Exceptions) + Chapter 7 (Comprehensions + scoping). - Weekend: **File 04** Chapter 8 (Functions + closures) + Chapter 12 (Data structures + Big-O).

Week 3 — Your track + tail breadth - Days 11–13: **File 04** Chapters 9 (Typing + Protocol), 10 (Modules + packaging), 11 (Strings + bytes + regex). Skim if non-relevant; deep-read if your team writes in this surface. - Days 14–15: Open your track file. Read the first half of **File 06** (Data) or **File 07** (Backend). - Weekend: Finish the back half of your track file.

Week 4 — Intel, scenarios, dress rehearsal - Day 22: **File 03 What Changed in 2026**. This is your talk-track material when the interviewer asks "what's new in Python?" or "have you used the 3.12 type-parameter syntax?" Read it like a magazine, mark the bits you can speak to. - Days 23–24: **File 02 Top 200** — read like a frequency map. Compare against what you've already covered in File 04 and your track file. The B-tier is where new ground lives. - Days 25–26: Mock-pass — pull 15 random Qs from **File 04** + 10 from your track + 5 traps from **File 05**. Time yourself 90 seconds each, speak out loud. Note which Qs you couldn't finish. - Days 27–28: Re-attempt all gaps. Re-pass any chapter where 40%+ of Qs felt shaky. - Day 29: **File 08** Cheatsheet — read until you can recite the 8 traps and the time-complexity table without looking. - Day 30 (interview eve): **File 08** once more. Skim **File 01**. Stop studying. Sleep.

Optional throughout: File 09 Python-Role Salary Map the day before negotiation, never before the technical round. The salary data is for after the offer, not before the test.

How To Use File 08 Cheatsheet On Interview Day

File 08 Cheatsheet is 1 dense page. Treat it like an exam aid, not study material.

- **Print it on A4, fold in half, keep in your bag.** Glance at it in the auto, in the lobby, in the waiting room. Do not carry it into the panel room — neither offline nor in your laptop's foreground tab.
- **Online interview:** pin the PDF on your second monitor or print and tape next to your webcam, angled so you can glance without losing eye contact. Interviewers don't mind; many of them have notes too.

- **5 minutes before the round:** read the cheatsheet once, slowly, top to bottom. The two highest-leverage sections are the **8 traps row** (one of these categories shows up more often than not) and the **time-complexity table** (every DSA-in-Python question hinges on `list.append` being $O(1)$ amortised, `list.insert(0, x)` being $O(n)$, `dict.get` being $O(1)$, `in` on a list being $O(n)$ — get these wrong and you sound junior).
 - **During the round:** if the interviewer asks "what's the time complexity of `bisect.insort`?" you should answer in 2 seconds without looking. The cheatsheet is for confidence-of-recall, not lookup.
 - **The pandas quick-reference (page bottom-right):** only matters if you're on the Data track. Memorise the `groupby().agg()` multi-column syntax and the `merge` `validate=` flags. Round 1 will ask one of these by name.
-

How The Drive Folder Works

This is a Drive link, not a ZIP. **Bookmark the folder, don't redownload PDFs.** Lifetime updates means the same folder refreshes — when we ship the next edition with new traps, new scenarios, and 2026-Q3 currency, the files you bookmarked update in place. If you saved local copies to your laptop, you'll be reading stale Qs six months from now while everyone else is on the current edition.

Open the folder. Find your row from the diagnostic above. Open that file. Go.