

Python-Role Salary Map — India, 2026

Anchor your number. Negotiate from data, not gut feel.

The ranges here are aggregated from public salary signals (Levels.fyi India, AmbitionBox, Glassdoor India, offer screenshots circulating in dev Telegram groups and r/developersIndia) plus pattern-matching from candidate reports seen over the last 12 months. Treat them as anchoring, not gospel. Your number depends on your specific company, team, the niche skill you bring, current CTC, urgency of the hiring slot, and the mood of the hiring manager that quarter. If your number lands in the middle of the right range, you're priced fairly. If you're 30% below midpoint, you should be having a different conversation — with your manager or with the market.

How to read this map

All numbers are **total CTC in LPA** (Lakhs Per Annum) — base + variable + estimated ESOP/stock value at offer, not just base. India compensation is almost always quoted as CTC in offer letters, so we match that convention.

Variable pay assumption per tier:

- **Service MNCs:** variable is 5–10% of CTC, almost always paid in full unless the company missed quarterly targets.
- **Indian Product cos:** variable is 10–20% of CTC, payout linked to individual rating + company performance. Plan for 80–90% in a normal year.
- **GCCs:** variable is 10–15%, usually paid out close to 100% — these orgs treat variable as deferred fixed.
- **Startups:** variable is small (0–10%); the rest is ESOP. ESOPs in Indian startups are often valued at zero on exit — treat them as bonus optionality, not real pay.

When a range says "₹8L–14L", read it as: 25th percentile around 8, median around 11, 75th percentile around 14. Hitting the top of the range needs either a niche skill (Section: "What's growing in 2026") or a track record (multi-year delivery on a P0 system, hard numbers, on-call ownership).

How Python roles map to comp

Hiring titles in India are messy. The same job is called four different things across four companies. Here is how we are bucketing Python roles for this map.

Data Analyst — SQL plus pandas for ad-hoc analysis, plus a reporting tool (Tableau / Looker / PowerBI / Metabase). Spends the day answering questions from product, marketing, ops, finance. Python here is the second language after SQL, used mostly for cleaning, joins pandas can do better than SQL, and quick scripts.

Data Scientist — Statistical modelling, experimentation, A/B test analysis, and increasingly ML work. Python is the primary tool: pandas + numpy + scikit-learn + statsmodels, with PyTorch / TensorFlow for the ML-leaning subset. Title often means very different things across companies — "data scientist" at a small startup is what an analytics engineer does elsewhere; "data scientist" at Flipkart or Swiggy is closer to applied ML.

Data Engineer — Pipelines. Airflow / Dagster / Prefect for orchestration, PySpark / Spark / dbt for transformation, Snowflake / Databricks / BigQuery as the warehouse, Kafka / Kinesis if streaming. Python is the glue and the transformation language. Your customer is the data scientist, the analyst, the ML engineer.

ML Engineer — Productionising models. Python from end to end: PyTorch / TensorFlow / scikit-learn for training, MLflow / Weights & Biases for tracking, FastAPI / Triton / TorchServe for serving, Kubernetes + Airflow for orchestration. The fastest-growing comp band in the Python pool, especially with GenAI exposure (LangChain, vector DBs, embedding pipelines, RAG systems).

Backend Dev (Python) — Django / Flask / FastAPI engineer writing application services. Owns transactional schemas, query performance for product features, REST or GraphQL contracts, background jobs (Celery / RQ / Dramatiq). Not a data engineer — primary deliverable is the API or the service.

Full-stack Dev (Python on backend) — Backend Dev plus a frontend skill (React / Vue / Next.js). Common at smaller startups where headcount is tight; less common at product cos where backend and frontend are separate tracks. Comp tracks closely to Backend Dev with a small premium when the frontend skill is real.

SDE generalist — The "writes Python because it's the team's default" SDE at a product co or GCC. Comp is set by SDE band (SDE-1 / SDE-2 / SDE-3 / Staff), not by language. Python here is incidental — the candidate could equally be writing Go or Java. Bench role at Service MNCs maps here too.

Note: in India 2026, Data Scientist comp tracks at or slightly below Data Engineer in most tiers — DE specialization commands more given the rise of ML platforms. ML Engineer is where the +10-17% premium lives.

Salary ranges by company tier

All numbers are total CTC in LPA. Read as 25th–75th percentile of public reports.

Tier 1 — Service MNCs (TCS, Infosys, Wipro, HCL, Accenture, Cognizant, Capgemini, LTIMindtree)

Lower base, less variance, predictable raise cycles (8–12% YoY on average, with promotion bumps of 25–40%). The offer-letter number is approximately what you'll take home — no hidden ESOP gravy.

Role	Fresher (0–1y)	1–3y	3–7y	7–12y	12y+
Data Analyst	3.5–5	5–8	8–14	13–22	20–32
Data Scientist	4–6	6–10	10–18	17–28	26–42
Data Engineer	4–6	6.5–11	11–19	18–30	28–45
ML Engineer	4.5–6.5	7–12	12–22	20–34	30–50
Backend Dev (Python)	3.5–5.5	5.5–10	10–17	16–26	22–35
Full-stack Dev	3.5–5.5	5.5–10	10–17	16–26	22–35
SDE generalist	3.5–5	5–9	9–15	14–24	22–36

Notes: - ML Engineer and Data Engineer roles at Service MNCs typically require client deployment at a tier-1 GCC client (JPMC, Walmart Labs, Deutsche Bank). Numbers above assume that. Pure-bench ML / DE roles are rare and pay closer to Backend Dev. - Backend Dev and Full-stack at this tier are functionally indistinguishable in offer terms — both are "developer" band with the same salary structure. - Service MNC freshers regardless of role start at the standard ₹3.5–4 LPA band; the spread shows up after the first lateral promotion (typically year 2–3).

Tier 2 — Indian Product Cos (Flipkart, Razorpay, Zerodha, Swiggy, Paytm, PhonePe, Freshworks, Zoho, Postman, Cred, Meesho, Groww, Nykaa, Urban Company)

Higher base plus a meaningful ESOP component. Roughly 3–3.5× Service MNC numbers for the same experience band.

Role	Fresher (0–1y)	1–3y	3–7y	7–12y	12y+
Data Analyst	9–15	15–26	24–42	38–60	55–85
Data Scientist	14–22	22–38	38–62	58–95	85–140
Data Engineer	14–24	22–40	38–65	60–100	90–150
ML Engineer	16–26	26–45	44–75	70–120	100–170
Backend Dev (Python)	12–22	20–38	36–60	55–95	80–130
Full-stack Dev	12–22	20–38	36–60	55–95	80–130
SDE generalist	14–24	22–42	40–65	60–100	85–140

Important caveat: 20–40% of these numbers may be ESOP. ESOPs in Indian private companies often value at zero on exit — the company didn't IPO, secondary market refused to clear, or the strike price was set high enough that the spread on a tender offer was thin. Treat ESOP as bonus optionality. Negotiate hard on base + cash variable. If the recruiter is leaning on the ESOP number to justify a low base, push back.

Zoho is a known outlier within this tier — its comp structure runs roughly 20–30% lower than the rest of the cohort because it doesn't offer ESOP and refuses to compete on cash. Adjust accordingly if you're benchmarking against a Zoho offer.

Tier 3 — GCCs (Microsoft India, Google India, Amazon India, JP Morgan, Goldman Sachs, Morgan Stanley, Walmart Labs, Target, Lowe's, Wells Fargo, AmEx, Bank of America, Citi, Deutsche Bank, Atlassian, Salesforce India)

Mid-to-high base, lower variable, very stable. ESOPs are real (US-listed parent). Often beat product cos at senior levels (8+ years) because Indian product cos cap senior IC growth somewhere around principal-engineer-ish bands.

Role	Fresher (0–1y)	1–3y	3–7y	7–12y	12y+
Data Analyst	10–18	18–30	28–48	45–75	70–110
Data Scientist	16–26	25–42	42–70	68–115	100–170
Data Engineer	15–28	26–45	45–75	70–125	105–180
ML Engineer	18–30	28–50	50–85	80–140	120–200
Backend Dev (Python)	14–25	24–42	42–70	65–110	90–150
Full-stack Dev	14–25	24–42	42–70	65–110	90–150
SDE generalist	16–28	26–48	48–80	75–130	105–175

Notes: - The top of the 12y+ range (₹1.5–2 Cr) is Big Tech and tier-1 IBs — Google, Microsoft, Amazon, Atlassian, JPMC strats, Goldman tech. Standard GCCs (Walmart, Target, Wells, Citi, Lowe's) typically cap around ₹70–110 LPA for principal-level Python ICs. - RSU vesting at US-listed parents is the real differentiator vs Indian product cos. Microsoft and Google grant new RSU refreshes annually; standard GCCs grant at hire + at promotion. - ML Engineer numbers at the top tier (Microsoft, Google) routinely exceed the table here when the candidate has applied research depth — those are recruiter-driven offers, not band-bound.

Tier 4 — Startup Series A–C (\$5M–\$50M raised; 30–250 headcount)

Base is lower than product cos. ESOP is higher as a percentage of CTC — sometimes meaningful (think early Razorpay, Cred, Zerodha hires), often not. Honest version: most Series A/B startup ESOPs don't pay out. Plan accordingly.

Role	Fresher (0–1y)	1–3y	3–7y	7–12y	12y+
Data Analyst	7–13	13–22	22–38	32–55	48–75
Data Scientist	10–18	18–32	32–55	48–80	70–110
Data Engineer	10–18	18–32	32–55	50–85	75–120
ML Engineer	12–20	20–35	35–60	55–95	80–135
Backend Dev (Python)	8–16	16–28	28–48	42–72	60–95
Full-stack Dev	8–16	16–28	28–48	42–72	60–95
SDE generalist	9–17	17–30	30–52	48–80	68–105

Notes: - At Series A/B you typically join a 2–10 person engineering or data team and own a vertical. The compensation discount vs product cos (roughly 15–25%) is supposed to be made up by ESOP and learning velocity. ESOP rarely lands; learning velocity often does. Decide which one matters more to you.

- Series C startups tracking toward IPO (think Razorpay 2022, Zerodha at any point) sometimes pay closer to product-co rates with real ESOP optionality. Treat the C-stage end of this tier as a hybrid.

Tier 5 — Late-stage / unicorn (\$50M+ raised; 250+ headcount; pre-IPO band)

Comp moves closer to product-co numbers because they're competing for the same talent pool. ESOP is more credible than Series A/B (some genuine secondary liquidity, occasional tender offers), still risky. Distinct band from product cos because the path to liquidity is longer and the culture is more chaotic.

Role	Fresher (0–1y)	1–3y	3–7y	7–12y	12y+
Data Analyst	10–18	18–30	28–48	45–72	65–100
Data Scientist	15–24	24–42	42–68	62–100	90–145
Data Engineer	15–26	25–44	42–70	65–110	95–160
ML Engineer	17–28	28–48	46–80	75–125	105–180
Backend Dev (Python)	13–24	22–42	40–65	60–100	85–140
Full-stack Dev	13–24	22–42	40–65	60–100	85–140
SDE generalist	15–26	25–46	44–72	65–110	90–150

Notes: - Late-stage examples in 2026: Razorpay, Cred, Meesho, Groww, Postman, Browserstack, Innovaccer, Sarvam (growing fast), Krutrim. Comp depends heavily on the last funding round's date and valuation. A company that raised in 2021 at peak valuations is paying lower base + higher (paper-rich, exit-poor) ESOP than one that raised at 2024 valuations. - This tier rewards niche skill stacking more than tier 4 — they have the budget and they're racing competitors for the same Snowflake / streaming / GenAI engineers.

What moves you up a tier

The most common career mistake we see in candidate reports: assuming your next promotion will move you to the next tier. It almost never does. Tier jumps come from switching, not from internal promotion. Here is the pattern.

Inside a tier, you move bands. Across tiers, you switch.

A senior data engineer at TCS doesn't get promoted to an equivalent Razorpay band. They get promoted to a slightly-better-paid senior data engineer at TCS. The Razorpay version requires interviewing out — usually after the senior promotion, when the resume reads better.

The 2–3 year switch cycle is the rhythm most data and Python engineers actually run on. Here is what it looks like for a typical mid-career trajectory:

- Year 0: join Service MNC as fresher, ₹4 LPA.
- Year 2: internal promotion to senior associate, ₹6.5 LPA.
- Year 2.5: switch to a Series A/B startup or smaller product co for ₹12 LPA — first tier jump.
- Year 4–5: internal growth at startup brings you to ₹16–18 LPA. Maybe a mid-stage promotion to "senior engineer".
- Year 5: switch to a product co or GCC for ₹28–35 LPA — second tier jump.
- Year 7–8: internal growth + skill specialization brings you to ₹45–55 LPA at the GCC.
- Year 8: switch (or counter-offer) for the next tier band.

Three internal promotions plus three external switches in eight years is a clean median for the people who outpace the LinkedIn average.

What moves you up inside a tier

- **Take a lead role even without the title.** Onboard juniors, run the team standup, own on-call rotation, lead the design doc on a P0 system. Comp bumps follow.
- **Niche skill that the team needs and nobody else has.** PySpark at scale, Airflow architecture, distributed training with PyTorch, ML serving at production volume, real-time streaming pipelines. The premium is real and stacks (see "What's growing in 2026").

- **Cross-team mobility.** Move from backend to data infra, from analyst to data engineer, from data engineer to ML engineer. Each pivot is a documented expansion of scope and triggers a band re-evaluation.
- **Ship one P0 system end-to-end with hard numbers.** "I cut Snowflake spend from \$14K/month to \$6K by reclustering three tables and rewriting four cron queries" — specific dollar amounts beat vague claims every time.
- **Document outside your job.** A blog with 4 well-researched posts, a conference talk, an OSS contribution that landed in a project people use. Signal of taste. Promotion-cycle ammunition.

When to negotiate

- **At offer stage, always.** Even a polite single email pushing for 8–12% more works 50% of the time. Downside is zero; recruiters expect a counter.
- **At performance-review cycle, if your manager has just delivered a high rating.** The "I'm grateful for the rating; can we discuss the band move?" conversation works when timed against the rating, not 3 months later.
- **When a competitor reaches out and you have a real conversation.** Bring the offer-in-hand to your current manager; product cos and GCCs match more often than you think when the loss is concrete.
- **Don't negotiate during the first 60 days.** You don't have the political capital yet. Land first; renegotiate at the 6-month review or the next switch.

When not to negotiate

- Recruiter outreach with no offer. You have nothing to anchor against.
- Inside Service MNC promotion cycles — bands are locked, the manager has no discretion. Wait for the switch.
- During probation. The answer is no, and you've burned a small amount of trust asking.

What's growing in 2026

The Python comp landscape has bifurcated. ML, data platform, and GenAI engineering roles are running 15–30% premiums over comparable backend roles at the same tier. Pure-Python-CRUD backend at Service MNCs is being commoditized. Here is the breakdown.

The AI-adjacent premium (10–30% over comparable backend)

Across all five tiers, roles that touch ML pipelines, vector databases, embedding stores, LLM serving, RAG architectures, or model monitoring are commanding a premium. The premium scales with how production-real the work is:

- Data engineer with PySpark + Airflow but no ML touchpoints: baseline.

- Data engineer who also runs the feature store, owns embedding pipelines, or maintains vector DB infra: +10–15%.
- ML engineer who productionises models end-to-end (training → serving → monitoring): +10–17% over the baseline DE band.
- ML engineer with GenAI in production (LangChain / LlamaIndex / vector DB / RAG running at customer scale): +20–30%, with the top of the range capped only by the budget of the hiring company.

GenAI specifically is the loudest niche in 2026 hiring. Most postings asking for "LangChain + vector DBs + RAG" are 18-month-old startups burning runway. The ones that are real (Sarvam, Krutrim, Innovaccer, Cred's AI team, Razorpay's risk-AI team, the Microsoft / Google India ML orgs) are paying genuine premiums and competing hard for senior ML engineers.

Python-as-glue is growing fastest

The fastest-growing comp bands across all tiers are the roles where Python is the connective tissue:

- **Data infrastructure engineer** — Airflow + Spark + Snowflake + dbt + maybe Iceberg. The plumbing job. Premium because the workload doesn't go away and the supply is thin.
- **MLOps engineer** — MLflow + Kubeflow + KServe / Triton + monitoring + drift detection. Newer title, smaller talent pool, real budget.
- **Feature store engineer** — Feast / Tecton / in-house store + offline / online sync + lineage. Niche, growing.
- **Vector DB / embedding pipeline engineer** — Pinecone / Weaviate / Qdrant / pgvector + embedding model serving + retrieval evaluation. Emergent in 2024, mainstream by mid-2026.

Backend Python is being commoditized at Service MNCs

The plain Django / Flask CRUD job at a Service MNC is the slowest-growing band in this map. Service MNC Backend Python at the 3–7y band has moved from ₹10–14 LPA in 2022 to ₹10–17 LPA in 2026 — barely an inflation-adjusted hold. The supply of mid-level Backend Python developers far exceeds the demand for them at this tier.

Product cos and GCCs are still paying competitively for Backend Python because the platform engineering work (event-driven architecture, observability, async pipelines, API gateway layers) requires depth that the Service MNC bench doesn't typically build. If you're a Backend Python dev at a Service MNC and you want a real raise, the answer is a switch — and the switch is easier with one niche layer (async patterns at scale, FastAPI in production, a streaming systems exposure) than with another year of generic experience.

Niche skill premiums (additive to the base CTC range)

These stack on top of the role × tier × experience number above. Two premiums don't add linearly. If you have three of these, you're at the top of the range, full stop.

Skill	Premium	Applies to
PyTorch + production model serving (Triton / TorchServe / Ray Serve)	+15–25%	ML Engineer
LangChain + vector DB + RAG in production at customer scale	+20–30%	ML Engineer, Backend Dev
PySpark at production scale (Databricks / EMR, terabyte-day jobs)	+10–20%	Data Engineer
Airflow architecture (>500 DAGs, custom operators, CI/CD pipelines)	+10–15%	Data Engineer
Real-time / streaming Python (Kafka + Flink Python API or Spark Streaming, end-to-end ownership)	+20–30%	Data Engineer, ML Engineer
dbt + Python integrations (dbt-python models, custom adapters)	+10–15%	Data Engineer, Analytics Engineer
FastAPI at production scale (high-throughput async services, dependency injection in real systems)	+10–15%	Backend Dev
Django at scale (multi-tenant SaaS, sharded ORM, custom middleware patterns)	+10–15%	Backend Dev
MLflow / Weights & Biases / Kubeflow in production	+10–20%	ML Engineer
Feast / Tecton feature store ownership	+15–25%	ML Engineer, Data Engineer — scarcest
Vector DB ownership (Pinecone / Weaviate / pgvector at production load)	+15–25%	ML Engineer — newest premium
GenAI evaluation expertise (retrieval evals, RAG benchmarking, embedding model comparison)	+15–20%	ML Engineer

Stacking note: "PySpark + Airflow + Snowflake" stacks at roughly +25–30% total, not +40–60%. The market recognises the combined picture but doesn't pay double for overlapping skills. "LangChain + vector DB + RAG production" treats the three as one premium worth +20–30%, not three separate ones.

Bench rates — the Service MNC reality

The version that doesn't show up on LinkedIn. If you're a Python developer at a Service MNC who has just rolled off a project, here are the real numbers.

Monthly bench stipend by tier

- **Junior bench (0–3 years):** ₹25K–₹45K per month (which is your standard monthly salary; bench just means you're getting paid without billable work). At fresher band that's roughly ₹3–5.5 LPA pro-rated.
- **Mid bench (3–7 years):** ₹50K–₹95K per month. ₹6–11 LPA pro-rated. Some Service MNCs cut the variable component during bench; check your offer letter language carefully.
- **Senior bench (7y+):** ₹95K–₹1.5L per month. ₹11–18 LPA pro-rated. Senior bench is rare because the company has more incentive to find you a deployment.

The actual cash hitting your account doesn't drop on bench. What drops is the variable pay (which is usually a sub-10% portion at this tier anyway) and your trajectory. A bench longer than 3 months starts to show in your appraisal cycle — managers code "low utilization" against your rating, which compounds at the next raise cycle.

How long benches typically last

- 60% of benches end within 4 weeks. The talent supply team has visibility into upcoming demands and pre-routes you.
- 25% extend to 4–10 weeks. You'll be told "we're working on it" and pulled into a couple of internal training tracks to keep utilization on paper. The training is sometimes useful, mostly performative.
- 10% extend to 10–20 weeks. This is the "we don't have a slot for your skill profile" range. You'll start hearing about reskilling tracks (cloud certifications, framework switches).
- 5% extend beyond 20 weeks. At this point you're functionally redundant. Most companies will offer a voluntary exit package (1–3 months of severance) before this happens.

What to do during a bench

The brutal truth: most engineers waste a 6-week bench scrolling LinkedIn and binge-watching tutorials. The ones who use it well do one of three things.

1. **Switch your skill profile inside the company.** If you've been on a Java + Oracle stack and you can credibly add Python + Snowflake or Python + AWS, the talent supply team has a wider deployment surface for you. The internal Snowflake certification + a side project on a customer-like dataset is a 4-week investment that genuinely changes your deployment options.
2. **Get a cloud certification that the market actually pays for.** AWS Solutions Architect Associate or GCP Professional Data Engineer are the two that move your CTC at the next switch by ₹2–4 LPA. They're 80 hours of focused study and ₹15K out of pocket. Bench is the only time you'll have a contiguous 80 hours.

3. **Interview out.** A 6-week bench is roughly when most Service MNC Python developers should start the switching conversation anyway. Product cos and GCCs pay 3–3.5× the Service MNC band. The bench period coincides with appraisal cycles, recruiter outreach is at a seasonal peak, and you have time to prep.

The LinkedIn theater vs the reality

LinkedIn will tell you bench is an opportunity. The bench in practice is a slow squeeze on your motivation and your skill currency. The Python skills you used on the last project (some Django version, some pandas patterns) stop being current. Your reflexes dull. The interviewer at your next interview can tell within 10 minutes whether you've been writing production code in the last quarter or you've been on bench.

The honest play: a Service MNC bench longer than 8 weeks is a signal to interview out. The company is telling you it can't keep you billable, which means the next raise cycle will reflect that. By the time the talent supply team finds you a new project, the recruiter at Razorpay or Walmart Labs will have already moved your CTC by ₹4–8 LPA. Take the recruiter call.

One ugly truth about Python at Service MNCs

Service MNCs over-hire Python developers relative to demand because Python is the cheapest tech-track to ramp freshers into. The result: more Python developers on bench at any given moment than Java or DotNet. If you joined a Service MNC as a Python developer in 2023 or 2024, you're probably reading this on a bench. The good news: the market outside your company is paying real money for Python skills that include Snowflake, Airflow, FastAPI at scale, or any ML / data platform exposure. The bench is the window for the skill switch that funds your exit. Use it.

One-line salary check

If your current CTC is **30%+ below the midpoint** for your role × experience × company tier in this map, you're underpaid. Either fix that conversation internally in the next 6 months, or start interviewing.

If you're **20%+ above the top end** of your band, you're doing something most people in the market can't do. Protect it, document it, don't switch for less than a 50% jump. Anything less and you're trading certainty for a story the new company hasn't earned the right to tell yet.