

00 — Start Here: The 60-Second Routing Decision

You just unzipped 9 PDFs. You're staring at filenames. Your interview is somewhere between three days and three months away, and you're not sure which file to open first.

Don't open any of them yet.

Spend the next 60 seconds on this file instead. It exists for exactly one reason — to tell you which of the other eight PDFs you should open right now, in what order, and what to skip. Read the diagnostic below, find your row, then close this file and go.

The 3-Question Diagnostic

Answer these three questions in your head. Then jump to the matching row in the table after Q3.

Q1. When is your interview?

- (A) This week — 3 to 7 days
- (B) Next 2 weeks — 8 to 14 days
- (C) This month — 15 to 30 days
- (D) 1 to 3 months out
- (E) No specific timeline — I'm building the skill

Q2. What type of company?

- (a) Service MNC — TCS, Infosys, Wipro, Accenture, Cognizant, HCL, LTIMindtree, Tech Mahindra
- (b) Product company or funded startup — Razorpay, Swiggy, Cred, Meesho, Zerodha, Flipkart, PhonePe, Zomato, Paytm
- (c) Data-role company — analytics-heavy product cos, D2C brands, ad-tech, fintech analytics teams
- (d) GCC of US firm — JP Morgan India, Walmart Labs, Wells Fargo, Goldman Sachs, Target, AmEx, Morgan Stanley, Deutsche Bank

Q3. What's your role?

- (i) Backend dev — Java / Python / Node + SQL

- (ii) Data engineer
 - (iii) Data analyst or analytics engineer
 - (iv) BI engineer
 - (v) DBA
-

Routing Table — Open These In This Order

Find the row that matches your **runway** (Q1) and your **company × role** combination (Q2 + Q3). File numbers in **bold** are non-negotiable opens; the rest are "if you have time after."

Your situation	Open first	Open next	Then	Skip for now
(A) This week + Service MNC + any role	01 Top 50, 07 Cheatsheet	02 Top 200 (skim joins + windows only)	03 What Changed 2026	04, 05, 06, 08
(A) This week + Product co + Backend/DE	01 Top 50, 07 Cheatsheet	05 Scenario Library (debug + optimize sections only)	06 Modern DB (your stack only)	02, 04 full, 08
(A) This week + Product co + Data analyst/BI	01 Top 50, 07 Cheatsheet	05 Scenario (what-will-this-return + index decisions)	06 Modern DB (BigQuery or Snowflake section)	02, 04 full, 08
(A) This week + GCC + Backend/DE	01 Top 50, 07 Cheatsheet	05 Scenario (transaction conflicts + optimize)	04 Chapter 5 (transactions + isolation) only	02, 03, 04 rest, 08
(A) This week + Data-role company + any role	01 Top 50, 07 Cheatsheet	05 Scenario (debug + data quality)	06 Modern DB (your warehouse)	02, 04 full, 08
(B) 2 weeks + Service MNC	01 Top 50, 02 Top 200	04 Chapters 1, 2, 3, 7	07 Cheatsheet day before	05 full, 06, 08
(B) 2 weeks + Product co / GCC + Backend/DE	01 Top 50, 02 Top 200	04 Chapters 2, 3, 4, 5	05 debug + optimize + index decisions, 07 Cheatsheet day before	04 Ch 6, 9, 08
(B) 2 weeks + Data analyst/BI/DE	01 Top 50, 02 Top 200	04 Chapters 2, 3, 7, 8	06 your warehouse, 07 Cheatsheet day before	04 Ch 5, 6, 9, 08
(C) 1 month + any combination	See 30-day plan below	—	—	—
(D) 1–3 months + any combination	04 Core Library end-to-end, then 05	06 Modern DB for your stack	02 Top 200 as final dress-rehearsal, 08	None

Your situation	Open first	Open next	Then	Skip for now
	Scenario Library end-to-end		Salary Map before negotiating	
(E) No timeline / learning	04 Core Library at 1 chapter per week	05 Scenario at 1 category per week	06 Modern DB, 03 What Changed 2026	None

One special-case note for DBAs — regardless of runway, after the row above, also open **04 Chapter 4 (indexes + optimization)** and **04 Chapter 5 (transactions + isolation + locking)** in full. These two chapters are your interview, full stop.

The 5-Day Cram Plan

If you have exactly 5 days, here's the hour-by-hour. Total commitment: ~3 hours/day on weekdays, ~5 hours on the weekend day. About 18 hours total.

Day 1 (3 hrs) — Top 50 + frame your weak spots - Hour 1: Read **01 Top 50** start to finish. Don't try to solve. Just read question + answer + framework. - Hour 2: Re-read **01 Top 50**. This time, mark every question where you couldn't have written the answer cold. - Hour 3: Open **07 Cheatsheet**. Read all three pages. Print it.

Day 2 (3 hrs) — Joins, windows, CTEs - Hour 1: From **04 Core Library**, do Chapter 2 (joins) — first 25 questions, answers covered. - Hour 2: From **04**, do Chapter 3 (windows + CTEs) — first 25 questions. - Hour 3: Re-attempt every question you got wrong from hours 1 and 2.

Day 3 (3 hrs) — Indexes, optimization, EXPLAIN reading - Hour 1: From **04**, do Chapter 4 (indexes + optimization) — first 30 questions. - Hour 2: From **05 Scenario Library**, do the "optimize this query" category — first 15 scenarios. - Hour 3: From **05**, do the "index decisions" category — first 15 scenarios.

Day 4 (3 hrs) — Transactions + your stack's modern DB - Hour 1: From **04**, Chapter 5 (transactions, isolation, locking) — first 20 questions. Memorise the four isolation-level tables in **07 Cheatsheet**. - Hour 2: Open **06 Modern DB Specifics** — read ONLY the section for the database your target company uses. Snowflake for analytics roles, Postgres 16 for most product cos, MySQL 8 for legacy stacks, BigQuery for ad-tech and Google-shop, SQL Server 2022 for finance GCCs. - Hour 3: From **05**, do "debug this query" — first 15 scenarios.

Day 5 (6 hrs, weekend) — Mock + cheatsheet drill - Hours 1–2: Re-read **01 Top 50**. You should now be able to answer 40+ of 50 cold. Mark the gaps. - Hours 3–4: Speed-run **02 Top 200** — read the question, give yourself 90 seconds to think, then read the answer. Don't write anything. - Hour 5: Re-read

07 Cheatsheet three times. Cover the page, recite from memory. - Hour 6: Re-attempt all questions you flagged across days 1–4.

What to skip on a 5-day plan: 02 Top 200 (full read), 04 Chapters 6, 8, 9, 10 (normalization, set ops, stored procs, modern features), 05 categories beyond debug/optimize/index decisions, 03 What Changed 2026, 08 Salary Map.

The 30-Day Deep Prep Plan

If you have a month, you can actually internalise this stuff instead of cramming. Plan: ~1 hour/day weekdays, 3 hours each weekend day. About 35 hours total.

Week 1 — Foundations and the 50 Most-Asked - Days 1–2: **01 Top 50** end to end. Solve every question on paper before reading the answer. - Days 3–5: **04 Chapter 1 (fundamentals) + Chapter 2 (joins)**. All 50+50 questions. - Weekend: **04 Chapter 3 (windows + CTEs)**. This single chapter is where most candidates lose the offer. Don't rush.

Week 2 — Performance and reliability - Days 6–8: **04 Chapter 4 (indexes + optimization)**. Slowest chapter — go deep. - Days 9–10: **04 Chapter 5 (transactions, isolation, locking)**. Memorise the isolation tables from **07**. - Weekend: **05 Scenario Library** — debug + optimize categories (all questions).

Week 3 — Breadth, schema thinking, and modern stack - Days 11–12: **04 Chapter 6 (normalization) + Chapter 7 (aggregations)**. - Days 13–14: **04 Chapter 8 (set ops) + Chapter 9 (stored procs)**. - Day 15: **04 Chapter 10 (modern features)**. - Weekend: **06 Modern DB Specifics** — your target warehouse end to end, plus skim the others.

Week 4 — Scenarios, intel, dress rehearsal - Days 16–18: **05 Scenario Library** — design schema + what-will-this-return + index decisions categories. - Days 19–20: **05** — transaction conflicts + data quality investigations. - Day 21: **02 Top 200 at Indian MNCs** — read like a magazine, mark the patterns. - Day 22: **03 What Changed in 2026** — talk-track material for "what's new?" questions. - Days 23–25: Re-attempt every question you flagged across the month. - Days 26–28: Mock-interview drill — pick 15 random questions from **04** + 10 from **05**, time yourself 4 minutes each, speak the answer out loud. - Day 29: **07 Cheatsheet** — re-read until you can recite the function patterns and index pitfalls without looking. - Day 30 (interview eve): Read **07 Cheatsheet** + skim **01 Top 50**. Sleep early. Eat properly.

Optional: read **08 Salary Map** the day before negotiation, not before the technical round.

How To Use The Cheatsheet On Interview Day

07 Cheatsheet is 1 page. Treat it like an exam aid, not study material.

- **Offline interview:** print it on A4, fold in half, keep in your bag. Glance at it in the waiting room or auto on the way. Do NOT carry it into the panel room.
 - **Online interview:** pin the PDF open on your second monitor, or print and tape next to your webcam — angled so you can glance without breaking eye contact. Interviewers don't mind; many of them have their own notes too.
 - **Just before the round:** read it once, top to bottom, slowly. The three isolation-level tables are the highest-leverage thing on the page — most candidates fumble READ COMMITTED vs REPEATABLE READ under pressure. Memorise the row order.
 - **During the round:** if the interviewer asks "which isolation level prevents phantom reads?" — you should answer in 2 seconds without consulting anything. The cheatsheet is for confidence-of-recall, not lookup.
-

What This Pack Does NOT Include

Be clear-eyed about scope so you don't go in expecting the wrong thing.

- **No live mock interview.** This is self-paced reading material. If you need a human asking you SQL questions in real time, book a peer-mock separately.
- **No resume review.** Question patterns and answer frameworks only. We don't touch your CV.
- **No career counseling or role-selection advice.** We don't help you decide between an offer at TCS vs a startup, or whether you should pivot from backend to data engineering.
- **No system design or coding rounds beyond SQL.** This is a SQL-only pack. DSA, LLD/HLD, behavioural — out of scope.
- **No company-specific question dumps.** We tell you the patterns Indian MNCs and product cos ask. We do not promise "the exact 5 questions Razorpay asked candidate X last Tuesday."
- **No video.** Everything is in PDFs you can read on a phone, tablet, or laptop. No subscription, no login.

What it IS: 1,400+ questions across 8 PDFs, frameworks for answering them, a cheatsheet for D-day, and the intel layer (Top 50, Top 200, 2026 shifts) so you walk in knowing what's coming.

Recommended First Action

Stop here. Go back to the diagnostic table. Find your row. Open that PDF now. The rest of the pack works only if you start with the right file — and the right file is almost never "the next one in numerical order."

Pick. Open. Go.