

# Data-Role Salary Map — India, 2026

*Anchor your number. Negotiate from data, not gut feel.*

---

The ranges in this file are aggregated from publicly reported numbers (Levels.fyi India, AmbitionBox, recent offer screenshots circulating in developer communities and Telegram groups) plus pattern-matching from candidate reports we have seen over the last 12 months. Use them as anchoring, not gospel. Your number depends on your specific company, team, the niche skill you bring, your current CTC, the urgency of the hiring need, and frankly the mood of the hiring manager that quarter. Treat this as a sanity check — if your number is in this map's middle, you're priced fairly; if you're 30% below midpoint, you should be having a different conversation with your manager or with the market.

---

## How to read this map

All numbers are **total CTC in LPA** (Lakhs Per Annum) — that's base + variable + estimated ESOP/stock value at offer, not just base. India compensation is almost always quoted as CTC in offer letters, so we match that convention.

Variable pay assumption per tier: - **Service MNCs**: variable is 5–10% of CTC, almost always paid in full unless the company missed quarterly targets. - **Product companies**: variable is 10–20% of CTC, payout linked to individual rating + company performance. Plan for 80–90% in a normal year. - **GCCs**: variable is 10–15%, usually paid out close to 100% — these orgs treat variable as deferred fixed. - **Startups**: variable is small (0–10%); the rest is ESOP. ESOPs in Indian startups are often valued at zero on exit — treat them as bonus optionality, not real pay.

When a range says "8–14 LPA", read it as: 25th percentile around 8, median around 11, 75th percentile around 14. Hitting the top of the range requires either a niche skill premium (see Section 5) or strong track record (Section 6).

---

## By role — what each role actually is

Hiring titles in India are messy. The same job is called four different things across four companies. Here's how we are bucketing roles in this map.

**Backend Developer (SQL-heavy)** — A Java / Python / Node / Go engineer who writes SQL daily for application workloads. Owns transactional schemas, query performance for product features, ORM debugging, and the occasional ETL written in application code. Not a data engineer — their primary deliverable is the API or the service, with SQL as a load-bearing tool.

**Data Analyst** — SQL plus a reporting tool (Tableau / Looker / PowerBI / Metabase), business-facing. Spends the day answering questions from product, marketing, ops, and finance. Builds dashboards, runs ad-hoc cohort analyses, writes weekly business reviews. Often the first data hire at a startup. Career path splits later into analytics engineering, data science, or product analytics.

**Analytics Engineer** — The role that exists because data analysts kept writing the same CTEs and dashboards kept breaking. Lives in dbt on top of Snowflake / BigQuery / Redshift. Owns the modeling layer — staging, marts, metrics definitions. Business-facing but more rigorous than a pure analyst. Newish role in India (~3 years old as a distinct title), still scarce, still pays a premium.

**Data Engineer** — Pipelines. Airflow / Dagster / Prefect for orchestration; Spark / Flink / dbt for transformation; Snowflake / Databricks / BigQuery as the warehouse; Kafka / Kinesis if streaming. Infra-facing role — your customer is the data scientist, the analyst, the ML engineer. You worry about late-arriving data, backfills, schema drift, and cost.

**BI Engineer / BI Developer** — Heavy reporting stack ownership: PowerBI / Tableau / Looker / Qlik / SSRS. The dashboard is the deliverable. In larger enterprises this is a distinct career track; in startups it often blurs into analytics engineering or data analyst. Premium when paired with strong SQL and semantic modeling skills (LookML, DAX, Tableau calculated fields at scale).

**DBA (Database Administrator)** — Postgres / Oracle / SQL Server / MySQL operations. Performance tuning, query plan reviews, index strategy, backup / restore, replication topology, HA, failover drills, version upgrades. A shrinking-but-still-paid role at large enterprises (banks, BFSI, telecom) and a quietly well-paid role at product companies that run their own databases instead of fully managed services.

---

## Salary ranges by company tier

All numbers are total CTC in LPA. Read as 25th–75th percentile of public reports.

### Tier 1 — Service MNCs (TCS, Infosys, Wipro, Accenture, Cognizant, HCLTech, Capgemini, LTIMindtree)

Lower base, less variance, predictable raise cycles (8–12% YoY on average, with promotion bumps of 25–40%). The number on the offer letter is what you'll roughly take home — there's no hidden ESOP gravy here.

| Role                    | 0–2 yrs | 2–5 yrs | 5–8 yrs | 8+ yrs |
|-------------------------|---------|---------|---------|--------|
| Backend Dev (SQL-heavy) | 3.5–5.5 | 6–10    | 11–18   | 18–30  |
| Data Analyst            | 3–5     | 5–9     | 10–16   | 15–25  |
| Analytics Engineer      | 4–6     | 7–11    | 12–18   | 18–28  |
| Data Engineer           | 4–6     | 7–13    | 14–22   | 22–35  |
| BI Engineer             | 3.5–5.5 | 6–11    | 11–18   | 17–28  |
| DBA                     | 4–6     | 7–12    | 13–22   | 22–35  |

Notes: DBA pays disproportionately well at this tier because banks and BFSI clients pay premium for Oracle / SQL Server depth. Service MNCs deployed at BFSI clients on long-running engagements often beat the top end here for true specialists.

## Tier 2 — Product Companies (Swiggy, Zomato, Razorpay, Cred, Flipkart, Myntra, PhonePe, Meesho, Zerodha, Groww, Nykaa, Urban Company, PolicyBazaar)

Higher base plus a significant ESOP component. Roughly 1.5–2.5× Service MNC numbers for the same experience band.

| Role                    | 0–2 yrs | 2–5 yrs | 5–8 yrs | 8+ yrs |
|-------------------------|---------|---------|---------|--------|
| Backend Dev (SQL-heavy) | 12–22   | 22–40   | 40–65   | 60–110 |
| Data Analyst            | 9–16    | 16–28   | 28–45   | 40–70  |
| Analytics Engineer      | 12–20   | 20–35   | 35–55   | 50–85  |
| Data Engineer           | 14–24   | 24–42   | 42–70   | 65–120 |
| BI Engineer             | 10–18   | 18–32   | 32–50   | 45–75  |
| DBA                     | 12–22   | 22–38   | 38–60   | 55–90  |

Important caveat: 20–40% of these numbers may be ESOP. ESOPs in Indian private companies are often valued at zero on exit — either because the company didn't IPO, the secondary market refused to clear them, or the strike price was set high enough that the spread on a tender offer was thin. Treat ESOP as bonus optionality. Negotiate hard on the base + RSU/cash component. If the recruiter is leaning on the ESOP number to justify a low base, push back.

### Tier 3 — GCCs of US/EU Firms (JP Morgan India, Goldman Sachs, Morgan Stanley, Wells Fargo, Walmart Labs, Target, AmEx, Bank of America, Citi, Deutsche Bank, Salesforce India, Microsoft India, Google India, Amazon India, Atlassian)

Mid-to-high base, lower variable, very stable. ESOPs are real (US-listed parent). Often beat product companies at senior levels (8+ yrs) because Indian product companies cap senior IC growth somewhere around principal-engineer-ish bands.

| Role                    | 0–2 yrs | 2–5 yrs | 5–8 yrs | 8+ yrs |
|-------------------------|---------|---------|---------|--------|
| Backend Dev (SQL-heavy) | 14–25   | 25–45   | 45–75   | 70–140 |
| Data Analyst            | 10–18   | 18–32   | 32–55   | 50–90  |
| Analytics Engineer      | 14–22   | 22–40   | 40–65   | 60–100 |
| Data Engineer           | 15–28   | 28–48   | 48–80   | 75–150 |
| BI Engineer             | 11–20   | 20–35   | 35–58   | 50–90  |
| DBA                     | 14–25   | 25–42   | 42–70   | 65–115 |

Notes: The top of the 8+ years range (₹1.2–1.5 Cr) is Big Tech and tier-1 IBs — Google, Amazon, Microsoft, Atlassian, JPMorgan tech, Goldman strats. Regular GCCs (Walmart, Target, Wells, Citi) typically cap around ₹60–90 LPA for principal-level data / backend ICs. RSU vesting at US-listed parents is the real differentiator vs Indian product companies.

### Tier 4 — Indian Startups (Series A / Series B+) (Series A: \$5–15M raised; Series B: \$15–50M)

Base is lower than product companies. ESOP is higher in percentage of CTC — sometimes meaningful (think early Razorpay, Cred, Zerodha hires), often not. The honest version: most Series A/B startup ESOPs don't pay out. Plan accordingly.

| Role                    | 0–2 yrs | 2–5 yrs | 5–8 yrs | 8+ yrs |
|-------------------------|---------|---------|---------|--------|
| Backend Dev (SQL-heavy) | 8–16    | 16–30   | 28–50   | 45–80  |
| Data Analyst            | 7–13    | 13–22   | 22–38   | 32–55  |
| Analytics Engineer      | 9–16    | 16–28   | 28–45   | 40–65  |
| Data Engineer           | 10–18   | 18–32   | 32–55   | 50–85  |
| BI Engineer             | 7–14    | 14–24   | 24–40   | 35–60  |
| DBA                     | 9–16    | 16–28   | 28–45   | 40–65  |

Notes: At Series A/B you typically join a 2–10 person data team and own a vertical. The compensation discount vs product cos (roughly 15–25%) is supposed to be made up by ESOP and learning velocity. The ESOP rarely lands; the learning velocity often does. Decide which one matters more to you.

### Tier 5 — Early-stage startups (pre-Series A / seed / pre-seed)

Base sometimes matches Service MNC numbers, sometimes worse. ESOP is the bet — usually 0.1–1% for the first 5–10 engineers, vesting over 4 years with a 1-year cliff.

| Role                    | 0–2 yrs | 2–5 yrs | 5–8 yrs | 8+ yrs |
|-------------------------|---------|---------|---------|--------|
| Backend Dev (SQL-heavy) | 5–12    | 12–22   | 22–40   | 35–60  |
| Data Analyst            | 4–10    | 10–18   | 18–30   | 25–45  |
| Data Engineer           | 6–14    | 14–25   | 25–42   | 38–65  |

(Analytics Engineer / BI / DBA roles are rare at pre-Series A — usually one generalist does all of it.)

High variance. The right early-stage bet at the right time pays better than 10 years at a GCC; the wrong one wastes 3 years and you re-enter the market with a CTC drop. Don't make this bet for the money — make it for the founder, the problem, and the equity terms.

## Niche skill premiums

These are additive percentages on top of the base CTC range above. Stack them carefully — two premiums don't usually add linearly. If you have three of these, you're at the top of the range, full stop.

| Skill                                                                                       | Premium | Applies to                           |
|---------------------------------------------------------------------------------------------|---------|--------------------------------------|
| Snowflake production experience (>1 yr, real workloads, real warehouses)                    | +15–25% | Data Eng, Analytics Eng              |
| BigQuery production experience (partitioning, clustering, BI Engine, cost-aware querying)   | +15–25% | Data Eng, Analytics Eng              |
| dbt expertise + analytics engineering practice (modeling, tests, exposures, semantic layer) | +15–20% | Analytics Eng, Data Analyst climbing |
| Spark + PySpark at production scale (Databricks / EMR, terabyte-day jobs)                   | +10–20% | Data Eng                             |
| Real-time / streaming (Kafka + Flink or Spark Streaming, end-to-end ownership)              | +20–30% | Data Eng — scarcest premium          |
| Postgres internals: replication, partitioning, performance tuning at scale (>500 GB)        | +10–20% | Backend Dev, DBA                     |
| Databricks production experience (Unity Catalog, Delta Live Tables, MLflow integration)     | +15–25% | Data Eng                             |
| Looker / LookML modeling (PDTs, derived tables, persistent caching strategy)                | +10–15% | Analytics Eng, BI Eng                |
| AWS Redshift / GCP BigQuery cost optimization with measurable \$ saved                      | +10–15% | Data Eng, Analytics Eng              |
| Airflow at scale (>500 DAGs, custom operators, CI/CD pipeline)                              | +10–15% | Data Eng                             |
| DAX / PowerBI semantic model design at enterprise scale                                     | +10–15% | BI Eng                               |
| Apache Iceberg / open table format production experience                                    | +15–25% | Data Eng — niche, growing            |

A note on stacking: "Snowflake + dbt + cost optimization" stacks as roughly +25–30% total, not +40–60%. The market is good at recognizing the combined picture but doesn't pay double for overlapping skills.

## What pushes you to the top of your range

Hiring managers fight harder for candidates who tick these. Each one is worth roughly 5–10% in your number.

- **Owned a P0 system end-to-end** — a revenue-critical pipeline, a customer-facing dashboard, the order-flow query that runs on every checkout. Not "contributed to". Owned: on-call, debug at 2 AM, design doc with your name on it.
- **Snowflake / BigQuery / Redshift cost optimization with hard numbers** — "I cut our warehouse spend from \$14K/month to \$6K/month by reclustering three tables and rewriting four cron queries." Specific dollar amounts beat vague claims every time.
- **Mentor / lead a small team without the title** — onboarded juniors, ran the data team standup, owned the on-call rotation. Recruiters can verify this in references.
- **Ran a migration** — Oracle to Postgres, on-prem warehouse to Snowflake, MySQL sharding, a Hadoop deprecation. Migrations require breadth, patience, and political capital. They pay.
- **Conference talk, OSS contribution, or written content showing domain depth** — even a moderate blog with 4 well-researched posts beats a quiet LinkedIn. Signal of taste.
- **Hired a junior who stayed 1+ years** — proves you can interview, onboard, and retain. Most engineers can't do all three.
- **Internal mobility track record** — switched teams 1–2 times, took on a new domain, moved from analyst to analytics engineer or backend to data eng. Signals adaptability.
- **You can explain trade-offs out loud** — "We picked Snowflake over BigQuery because our team already knew SQL Server semantics and our workload was bursty, not steady-state." Specificity wins interviews.
- **Quantifiable impact in numbers your CTO would care about** — latency reduced from 800 ms to 120 ms, weekly active users up 14%, ETL runtime from 6 hours to 40 minutes, an audit failure turned into a clean report.
- **A track record of saying "we shouldn't build this"** — pushed back on a feature, prevented an over-engineered solution, killed a project. Senior signal.

---

## What pushes you to the bottom of your range

Honest list. Recruiters and hiring managers see these patterns and adjust offers down.

- **Stayed in the same role for 5+ years with no scope expansion** — same team, same systems, no promotion, no lateral. Suggests stagnation.
- **No production ownership of any system** — "I worked on it" doesn't survive a 30-minute deep dive. If the interviewer can't ask you about an incident you debugged at 11 PM, you don't own anything.

- **Can't speak to specific numbers** — vague answers about revenue impact, dataset sizes, query latencies, or cost. If everything is "large" and "complex", the interviewer assumes neither.
  - **Only worked on internal tools, never customer-facing** — internal-tools-only engineers are perceived (sometimes unfairly) as not having shipped under real constraints.
  - **Can't articulate trade-offs** — "We used Snowflake because it's better" without a reason is the death sentence of a senior interview.
  - **Job-hopping every 12 months without a story** — switching is fine; switching with no narrative arc reads as opportunism.
  - **No mention of failure or hard-learned lessons** — "everything went smoothly" in 8 years means you weren't in the room when things went wrong.
- 

## Negotiation thresholds — when to push, when to accept

**Always negotiate:** Any offer that lands >10% below the midpoint of your role + experience + company-tier range in this map. Even a polite single email pushing for 8–12% more works 50% of the time. The downside is zero; recruiters expect a counter.

**Worth pushing on:** - **Signing bonus** — companies bend on this more than base because it doesn't recur in next year's salary budget. Especially true for product cos and GCCs. - **Notice-period buyout** — if you're servicing 60–90 days notice at your current company, ask the new employer to cover the buyout. They almost always will up to ₹2–5L. - **Variable-pay percentage** — push for a higher base and lower variable. Variable is a real risk: 80–90% payout in normal years means 0–60% payout in bad years. - **Joining date flexibility** — if you can join in 30 days, that itself is leverage. Use it. - **ESOP refresh cadence at product cos** — some give refresh grants every 2 years, some only at promotion. Ask.

**Hard to move:** - **Equity refresh in Service MNCs** — rare. Don't waste leverage here. - **Title** — often locked to internal job bands. You can push from "Senior Data Engineer" to "Lead Data Engineer" maybe once in a negotiation cycle; never twice. - **Location / remote** — varies dramatically by company. Ask in the first call, not at offer stage. - **Probation terms** — almost no company will waive a 6-month probation just because you asked.

**Walk-away signals:** - Offer is <70% of your researched range AND there is no credible 12-month path to growth. - Vague answers on raise cycle ("we review when we can", "depends on company performance"). - "We'll discuss compensation after probation" — never accept this. Discuss it now or walk. - Recruiter ghosts you for 2+ weeks after offer discussion — signals internal chaos. - Hiring manager refuses to discuss the team's last 3 attrition cases.

---



## The "joining bonus" play

Joining bonuses are often the easiest lever to move because the money comes from a different budget line — talent acquisition, not the salary budget. The new manager doesn't have to defend it as a recurring cost.

**Typical ranges** (one-time, paid in 1–2 instalments, almost always with a claw-back clause if you exit within 12–24 months):

- **Service MNC:** ₹50K–₹3L. Smaller, but doable. Push when switching from a competitor that they want you out of.
- **Product Company:** ₹2L–₹15L. Standard, especially for niche skills (Snowflake, dbt, real-time data).
- **GCC of a US firm:** ₹3L–₹20L. The highest, especially if they're replacing an attrition or filling a long-vacant slot. AmEx, JPMC, Goldman go up to ₹25L for hard-to-find senior data engineers.
- **Indian Startup (Series A/B):** ₹2L–₹8L. Some swap base for joining bonus to keep cap-table optics clean; treat as compressed base.

**Two practical asks** to layer on:

1. "Can I get the joining bonus in two parts — half at joining, half at the 12-month mark — so the claw-back stays simple?" Most recruiters say yes; this de-risks both sides.
2. "Can the notice-period buyout be paid directly to my current employer's invoice, not as a reimbursement?" Saves you the temporary cash crunch.

Always read the claw-back clause carefully. Standard is full repayment if you exit before 12 months, 50% if you exit between 12–24 months. Anything stricter than this — push back.

---

## YoY salary inflation note

A short honest framing: 2022–2023 was an anomaly. Indian product companies and GCCs ran a one-time correction, with 30–50% YoY jumps becoming common for mid-level data and backend engineers — driven by post-COVID hiring, the brief crypto/Web3 cash surge, and inflated US startup funding flowing into India. That window closed in mid-2023.

2024 and 2025 normalized to 10–15% YoY raises for internal promotions and roughly 25–40% jumps when switching jobs. 2026 is tracking similarly — early data from Q1 hiring suggests another normalized year with no broad correction in sight.

Two pockets are still seeing scarcity premiums in 2026: data engineers with production Snowflake + dbt + streaming exposure, and analytics engineers with strong LookML / semantic-modeling experience. Both groups are still seeing 35–50% job-switch jumps. Pure backend SQL roles and traditional DBA roles are flat-to-modest.

Don't plan your career assuming another 2022-style correction. Plan it assuming 12–15% internal growth and 25–35% switch premiums, with the niche-skill stacking being the only reliable way to outperform.

---

## One-line salary check

If your current CTC is **30%+ below the midpoint** for your role / experience / company-tier in this map, you are probably underpaid. Either fix that conversation internally in the next 6 months or start interviewing.

If you are **20%+ above the top end** of your band, you're doing something most people in the market can't do — protect that, document it, and don't switch for less than a 50% jump. Anything less and you're trading certainty for a story the new company hasn't earned the right to tell yet.